

Final Report

Group 1: Jiawen Chen, Brooke Felsheim, Elena Kharitonova, Xinjie Qian, and Jiarui Tang

4/30/2022”

Introduction

As human populations are rising across the world, so is the proportion of people that live in urban areas. Estimates from the *UN World Urbanization Prospects* (2018) indicate that over 4.2 billion people (55% of the global population) currently live in urban areas, and by 2050, an additional 2.5 billion people (68% of the global population) could be living in urban areas. More people living in urban areas calls for more space-, cost-, and energy-efficient systems of transportation as an alternative to cars. One such promising transportation alternative is the implementation of bicycle sharing programs.

Bicycle sharing programs are transportation schemes that allow individuals to rent bicycles on a short-term basis for either a set rate or for free. Most bicycle sharing programs have many computer-controlled bicycle rack “hubs” dispersed across a city that keep bikes locked and release them for use when a user enters the appropriate information/payment from a station or an app (Figure 1). A user can then ride the bike and return it to any other bicycle hub that is part of the same program. Many cities across the world have begun implementing bicycle sharing programs, including Chapel Hill, which has a Tar Heel Bikes sharing system¹. Systems like these provide convenient, inexpensive, and eco-friendly transportation options for individuals residing in a city.



Figure 1: A ‘hub’ of bicycles belonging to the Santander Cycles system in London. SOPA Images/Lightrocket via Getty Images

Successful implementations of bike sharing programs depend on proper management of these systems. It is important for a bike sharing program to provide a stable supply of rental bikes to its population so its users feel that they can rely on the system for their transportation needs. The analysis of bike sharing data allows for a better understanding of the demand of rental bikes in a city, which, in turn, can help inform a city about how to provide appropriate supplies of rental bikes for its population. For our project, we were interested in predicting the number of bikes rented within a given bike sharing system given information about weather, time of day, and date. We were also interested in assessing the most important variables for predicting bike rental counts. To answer these questions, we fit and evaluated a negative binomial generalized mixed model, a conditional inference tree, and a random forest model, using data from three publicly available bike sharing demand datasets.

The first dataset we use is a London bike sharing demand dataset downloaded from Kaggle² and provided by

¹<https://move.unc.edu/bike/bikeshare/>

²<https://www.kaggle.com/datasets/hmavrodiev/london-bike-sharing-dataset>

Transport for London³. This dataset contains hourly bike rental count observations over two years, from Jan 04 2015 - Jan 03 2017. The first full consecutive year of data was used as the training set in the analysis, and the second full consecutive year of data was held out as a test set.

The second dataset we use is a Seoul bike sharing demand dataset downloaded from the UCI Machine Learning Repository⁴ and provided by the Seoul Metropolitan Government⁵. This dataset contains hourly bike rental counts over one year, from Dec 1 2017 - Nov 30 2018. This was used as an independent test set. The third dataset we use is a Washington, D.C. bike sharing demand dataset downloaded from Kaggle⁶ and provided by Capital Bikeshare⁷. This dataset contains hourly bike rental counts over two years, from Jan 01 2011 - Dec 31 2012. This was used as an independent test set.

Each dataset contained hourly observations of bike rental count data. To simplify our analysis, we chunked the hourly data into three time blocks: [0:00 - 8:00), [8:00 - 16:00), and [16:00 - 24:00). To calculate the total bike count, we summed the hourly bike counts for each of the time blocks. Additionally, because temperature and humidity can be correlated with time of day, we chose to use the maximum and minimum daily temperature and humidity measurements for each 8-hour data point in order to avoid any issues of colinearity.

We created an R package named `bikeSharing` that includes methods for training and evaluating the negative binomial glmm and random forest models, as well as the processed source data from all three sets.

The below code can be used to install the package and load the package library. The zipped package source data, `bikeSharing_1.0.0.tar.gz` can be found in the Github repository for our project⁸.

```
if(!require("bikeSharing", quietly = TRUE))
  install.packages("package/bikeSharing_1.0.0.tar.gz", repos = NULL)
library(bikeSharing)
```

Once the package is loaded, the bike sharing data from all three sets becomes easily accessible through the variable names `london`, `seoul`, and `dc`. To directly load them into the R environment, one can simply run:

```
data("london", "seoul", "dc")
```

All three datasets contain bike rental count data as well as 11 additional weather-, time-, and date-related variables that were used as predictors in our models:

- Hour chunk (00:00 - 8:00, 8:00-16:00, 16:00-24:00)
- Weekend status (Yes/No)
- Holiday status (Yes/No)
- Season (Winter, Spring, Summer, Autumn)
- Minimum daily temperature (C)
- Maximum daily temperature (C)
- Minimum daily humidity (%)
- Maximum daily humidity (%)
- Wind speed (m/s)
- Presence of any rain or snow (Yes/No)
- Date (mm-dd)

The way that these variables are used within our models will be further described in the Methods section. For all of the analyses performed, the `london` dataset was divided into training and testing sets, where the training set contained all “Year 1” data (Jan 04 2015 - Jan 03 2016), and the testing set contained all “Year 2” data (Jan 04 2016 - Jan 03 2017).

```
london_train <- london[london$Year == "Year 1",]
london_test  <- london[london$Year == "Year 2",]
```

³<https://cycling.data.tfl.gov.uk>

⁴<https://archive.ics.uci.edu/ml/datasets/Seoul+Bike+Sharing+Demand>

⁵<https://data.seoul.go.kr>

⁶<https://www.kaggle.com/datasets/marklvl/bike-sharing-dataset>

⁷<https://ride.capitalbikeshare.com/system-data>

⁸<https://github.com/brookefelsheim/bios735-group1>

Methods

Negative Binomial Generalized Linear Mixed Model

Let y_{ij} be the number of bikes rented at hour chunk j of day i . Thus i ranges from 1 to 365, and j ranges from 1 to 3, corresponding to hour chunks [00:00 - 8:00), [8:00-16:00), [16:00-24:00). We assume that the number of bikes rented for a given hour chunk within a specific day follows a negative binomial distribution, so $Y_{ij} \sim NB(\mu_{ij}, \theta)$ using the Hilbe parameterization, so:

$$P(Y_{ij} = y_{ij} | \mu_{ij}, \theta) = \frac{\Gamma(y_{ij} + \theta)}{\Gamma(y_{ij} + 1)\Gamma(\theta)} \left(\frac{\theta}{\theta + \mu_{ij}} \right)^\theta \left(\frac{\mu_{ij}}{\theta + \mu_{ij}} \right)^{y_{ij}} \quad (1)$$

The mean of y_{ij} for each day i at hour chunk j is μ_{ij} , which we assume follows a negative binomial generalized linear mixed model with a random intercept. Thus it is determined from the following model:

$$\log(\mu_{ij}) = x_{ij}^T \beta + b_i \quad (2)$$

Let $x_{ij} = (1, \text{I}(\text{HourChunk}_{ij} = [8,16]), \text{I}(\text{HourChunk}_{ij} = [16,24]), \text{Weekend}_i, \text{Holiday}_i, \text{I}(\text{Season}_i = \text{Spring}), \text{I}(\text{Season}_i = \text{Summer}), \text{I}(\text{Season}_i = \text{Winter}), \text{Min_Temperature}_i, \text{Max_Temperature}_i, \text{Min_Humidity}_i, \text{Max_Humidity}_i, \text{Wind_Speed}_{ij}, \text{Rain_or_Snow}_{ij})^T$.

HourChunk_{ij} is a categorical variable corresponding to the hour chunk j of day i , with the reference hour chunk being [0:00, 8:00), Weekend_i is a binary variable with 1 if day i is a weekend, 0 if it is not. Holiday_i is a binary variable with 1 if day i is a holiday, 0 if it is not. Season_i is a categorical variable corresponding to which season day i is in, with the reference season being Autumn. The Min_Temperature_i and Max_Temperature_i are the minimum and maximum temperature of day i , respectively. Similarly, the Min_Humidity_i and Max_Humidity_i are the minimum and maximum humidity of day i , respectively. Wind_Speed_{ij} is the average wind speed of hour chunk j of day i . Rain_or_Snow_{ij} is a binary variable with 1 if during day i hour chunk j there is any rain or snow, 0 if there is not. Thus $\beta = (\beta_0, \beta_1, \beta_2, \beta_3, \beta_4, \beta_5, \beta_6, \beta_7, \beta_8, \beta_9, \beta_{10}, \beta_{11}, \beta_{12}, \beta_{13})'$.

In Equation (2), b_i is the unobserved random effect that the day i has on the number of rented bikes at any given hour chunk. It is assumed that $b_i \sim N(0, \sigma_\gamma^2)$.

Thus, of interest is the estimate θ , β , and σ_γ^2 . Equations 1 and 2 can be combined to obtain the following likelihood equation for a given day (i).

$$L(\theta, \beta, \sigma_\gamma^2 | \mathbf{y}, \mathbf{b}) = \prod_{i=1}^{365} \prod_{j=1}^3 p(y_{ij}) = \prod_{i=1}^{365} \left(\prod_{j=1}^3 f(y_{ij} | \mu_{ij}) \right) f(b_i | \sigma_\gamma^2) \quad (3)$$

Since the random effects b_i are unobservable, this means they must be integrated out of the above expression to obtain the likelihood, so:

$$\begin{aligned} L(\theta, \beta, \sigma_\gamma^2 | \mathbf{y}, \mathbf{b}) &= \prod_{i=1}^{365} \left[\int \left(\prod_{j=1}^3 f(y_{ij} | \mu_{ij}) \right) f(b_i | \sigma_\gamma^2) db_i \right] \\ &= \prod_{i=1}^{365} \left[\int \prod_{j=1}^3 \frac{\Gamma(y_{ij} + \theta)}{\Gamma(y_{ij} + 1)\Gamma(\theta)} \left(\frac{\theta}{\theta + e^{x_{ij}^T \beta + b_i}} \right)^\theta \left(\frac{e^{x_{ij}^T \beta + b_i}}{\theta + e^{x_{ij}^T \beta + b_i}} \right)^{y_{ij}} \frac{1}{\sqrt{2\pi\sigma_\gamma^2}} \exp\left(-\frac{b_i^2}{2\sigma_\gamma^2}\right) db_i \right] \end{aligned}$$

So thus, the log likelihood is found to be

$$l(\theta, \beta, \sigma_\gamma^2) = \sum_{i=1}^{365} \log \left[\int \prod_{j=1}^3 \frac{\Gamma(y_{ij} + \theta)}{\Gamma(y_{ij} + 1)\Gamma(\theta)} \left(\frac{\theta}{\theta + e^{x_{ij}^T \beta + b_i}} \right)^\theta \cdot \left(\frac{e^{x_{ij}^T \beta + b_i}}{\theta + e^{x_{ij}^T \beta + b_i}} \right)^{y_{ij}} \frac{1}{\sqrt{2\pi\sigma_\gamma^2}} \exp \left(-\frac{b_i^2}{2\sigma_\gamma^2} \right) db_i \right]$$

This log-likelihood will be maximized to obtain estimates for $\theta, \beta, \sigma_\gamma^2$ through an MCEM approach. Assuming that the b_i 's were known, we first define the complete data log likelihood as:

$$\log L_C(\theta, \beta, \sigma_\gamma^2 | \mathbf{y}, \mathbf{b}) = \log \left[\prod_{i=1}^{365} \left(\prod_{j=1}^3 f(y_{ij} | \mu_{ij}) f(b_i | \sigma_\gamma^2) \right) \right] = \sum_{i=1}^{365} \log \left(\prod_{j=1}^3 f(y_{ij} | \mu_{ij}) f(b_i | \sigma_\gamma^2) \right)$$

$$\text{So, } l_C(\theta, \beta, \sigma_\gamma^2 | \mathbf{y}, \mathbf{b}) = \sum_{i=1}^{365} \left(\sum_{j=1}^3 \log f(y_{ij} | \mu_{ij}) + \log f(b_i | \sigma_\gamma^2) \right)$$

Thus, for the MCEM algorithm, we will be maximizing the expectation of the complete data log likelihood, otherwise known as the Q-function at step t . So the Q-function is defined as the following:

$$Q(\theta, \beta, \sigma_\gamma^2 | y, \theta^{(t)}, \beta^{(t)}, \sigma_\gamma^{2(t)}) = E[l_C(\theta, \beta, \sigma_\gamma^2 | \mathbf{y}, \mathbf{b}) | y, \theta^{(t)}, \beta^{(t)}, \sigma_\gamma^{2(t)}]$$

$$= E \left[\sum_{i=1}^{365} \left(\sum_{j=1}^3 \log f(y_{ij} | \mu_{ij}^{(t)}) + \log f(b_i | \sigma_\gamma^{2(t)}) \right) \right]$$

$$= \sum_{i=1}^{365} \left[\int \left(\sum_{j=1}^3 \log f(y_{ij} | \mu_{ij}^{(t)}) + \log f(b_i | \sigma_\gamma^{2(t)}) f(b_i | y, \theta^{(t)}, \beta^{(t)}, \sigma_\gamma^{2(t)}) \right) db_i \right]$$

Where $\mu_{ij}^{(t)} = e^{x_{ij}^T \beta^{(t)} + b_i}$ and $f(b_i | y, \theta^{(t)}, \beta^{(t)}, \sigma_\gamma^{2(t)})$ is the posterior distribution of b_i given the observed data and the current parameter estimates. Let $Y_i = (y_{i1}, y_{i2}, y_{i3})^T$. The density function for b_i given $Y_i, \beta^{(t)}, \theta^{(t)}, \sigma_\gamma^{2(t)}$ is:

$$f(b_i | Y_i, \theta^{(t)}, \beta^{(t)}, \sigma_\gamma^{2(t)}) \propto f(Y_i | b_i, \theta^{(t)}, \beta^{(t)}, \sigma_\gamma^{2(t)}) \cdot f(b_i | \sigma_\gamma^{2(t)})$$

$$\propto \prod_{j=1}^3 \left(\frac{\theta^{(t)}}{\mu_{ij}^{(t)} + \theta^{(t)}} \right)^{\theta^{(t)}} \cdot \left(\frac{\mu_{ij}^{(t)}}{\mu_{ij}^{(t)} + \theta^{(t)}} \right)^{y_{ij}} \cdot \exp \left(-\frac{b_i^2}{2\sigma_\gamma^{2(t)}} \right)$$

$$\propto \sum_{j=1}^3 \left[(y_{ij} \log \mu_{ij}^{(t)} - (y_{ij} + \theta^{(t)}) \log(\mu_{ij}^{(t)} + \theta^{(t)})) - \frac{b_i^2}{2\sigma_\gamma^{2(t)}} \right] \quad (4)$$

If we could sample from this posterior distribution of b_i , we could approximate this integral using a montecarlo approach. So,

$$Q(\theta, \beta, \sigma_\gamma^2 | y, \theta^{(t)}, \beta^{(t)}, \sigma_\gamma^{2(t)}) = \frac{1}{M} \sum_{i=1}^{365} \sum_{k=1}^M \left(\sum_{j=1}^3 \log f(y_{ij} | \mu_{ijk}^{(t)}) + \log f(b_{ik} | \sigma_\gamma^{2(t)}) \right)$$

Where b_{ik} is one of the M samples of the posterior distribution b_i and $\mu_{ijk}^{(t)} = e^{x_{ij}^T \beta^{(t)} + b_{ik}}$.

To sample from this distribution posterior distribution of b_i , we employed the metropolis hastings algorithm with a random walk. So at set t , for each new b_i , we considered $b_i^* = b_i^{(t)} + \epsilon$ where $\epsilon \sim U(-\frac{1}{4}, \frac{1}{4})$. From this, the Metropolis Hastings ratio $R(b_i^{(t)}, b_i^*)$ was determined whether or not to accept this new b_i^* . For each new b_i^*

$$R(b_i^{(t)}, b_i^*) = \frac{f(y_{ij}|\theta, \beta, \sigma_\gamma^2, b_i^*)f(b_i^*|\sigma_\gamma^2)}{f(y_{ij}|\theta, \beta, \sigma_\gamma^2, b_i^{(t)})f(b_i^{(t)}|\sigma_\gamma^2)}$$

Thus, the new value for the $b_i^{(t+1)}$ is:

$$b_i^{(t+1)} = \begin{cases} b_i^* & \text{with probability } \min(R(b_i^{(t)}, b_i^*), 1) \\ b_i^{(t)} & \text{otherwise} \end{cases}$$

For the M-step, we will maximize the Q-function with respect to $\theta, \beta, \sigma_\gamma^2$. The Nelder-Mead method was used to optimize all of these parameters, this was because Nelder-Mead in the M-step does not require 1st or 2nd derivatives and it is robust. After this, the E-step and M-step are repeated for each iteration until the estimates for $\theta, \beta, \sigma_\gamma^2$ converge.

In our package **bikeSharing**, we have created a function **MCEM_algorithm** that will fit a Negative Binomial Generalized Linear Mixed Model to the data using the MCEM algorithm described above.

The **MCEM_algorithm** function takes as input **beta_initial**, which are initial guesses for the β vector. **theta_initial** which is the initial guess for the θ vector. **s2gamma_initial** which is the initial guess for the σ_γ^2 vector. **M** which is the number of posterior b_i 's to be sampled for the Monte Carlo integration. **burn.in** which is the number of posterior b_i 's to be thrown out to account for the time it takes for the chain to converge. **tol** which is the percent difference between the Q functions of two iteration to reach convergence. **maxit** which is the maximum number of iterations allowed. **data**, which is the data to which the negative binomial GLMM model should be fit to.

```
set.seed(1)
glmm_fit <- MCEM_algorithm( beta_initial = c(8.3, 1.5, 1.5, -0.25, -0.50, 0,
                                             0, -0.25, 0, 0, 0, 0, 0, -0.25),
                           theta_initial = 10,
                           s2gamma_initial = 0.2,
                           M = 1000,
                           burn.in = 200,
                           tol = 10^-4,
                           maxit = 100,
                           data = london_train
                           )
```

The output of the **MCEM_algorithm** function contains the list of **beta**, the final iteration of the fitted $\hat{\beta}$ vector. It contains **s2gamma** which is $\hat{\sigma}_\gamma^2$ at the final iteration. **theta** which is the final iteration of $\hat{\theta}$. **eps** which is the difference between this iteration and the last. **qfunction** which is the value of the Q function at the final iteration. **day_ranef** which contains all of the fitted b_i 's. **iter** is the iteration at which convergence was reached.

```
str(glmm_fit)

## List of 7
## $ beta      : num [1:14] 8.353 1.534 1.415 -0.337 -0.393 ...
## $ s2gamma    : num 0.0296
## $ theta     : num 18.4
## $ eps       : num 5.15e-05
```

```
## $ qfunction: num -9520
## $ day_ranef: num [1:365] 0.0653 -0.398 -0.5165 -0.2612 -0.0374 ...
## $ iter      : num 23
```




Figure 2: Values of theta in each iteration of the run of `MCEM_algorithm` on the London training set.

As we can see in Figure 2, the estimation of $\hat{\theta}$ converges quickly and from 9th iteration, it starts to float near the true θ .

The size of the coefficients of the binary and categorical variables can be compared directly to determine how important that variable is in determining bike count. As we can see from above, Hour Chunks [8,16) and Hour Chunks [16,24) are the covariates with the largest coefficients in magnitude, implying that they are very important in determining the number of bikes rented. As we can see, between 8:00 - 16:00 and 16:00 - 24:00, many more people rent bikes in London than between 0:00 - 8:00. This result is not surprising. Weekend and Holiday are the next binary variables with the largest coefficients in magnitude. Both holidays and weekends have a similar effect on bike count; less people rent bikes on holidays and weekends. Additionally, rain or snow causes people to rent less bikes. Winter has a larger effect on bike counts than spring or summer compared to autumn, with significantly less people renting bikes in winter than autumn. Less bikes are rented in spring than autumn although the effect is not as large as the effect of winter. Summer leads to an increase in bike rentals compared to autumn. Min and Max Temperature do not play as significant of a role in determining the number of bike counts, but this effect may be masked by the effects of the season. Humidity does not seem to play a large role in determining bike counts. Wind speed has a small negative coefficient, meaning that with increased wind speed, less people rent bikes.

Machine learning models

Next, we applied two machine learning methods to predict the bike count. The first model we used is a conditional inference tree. In traditional partitioning, all possible splits are investigated to find the best split, which results in overfitting and selection. The conditional inference tree embeds the partition step with a permutation test, thereby enabling this method to be robust to covariates of different scales. Furthermore, it is capable of stopping when no significant correlation exists between the covariates and the response (Hothorn, Hornik, and Zeileis 2006). We employed the `ctree` function in the `party` package to fit the conditional inference tree.

We visualized the three layer conditional inference tree in order to verify the relationship between bike count and variables (Figure 2). It is apparent that people tend to rent fewer bikes between midnight and 8 am. Maximum temperatures and humidity also affect bike rental rates. The splits that are employed in the conditional inference tree provide well reasoned explanations of the data structure, which means high accuracy predictions are made.

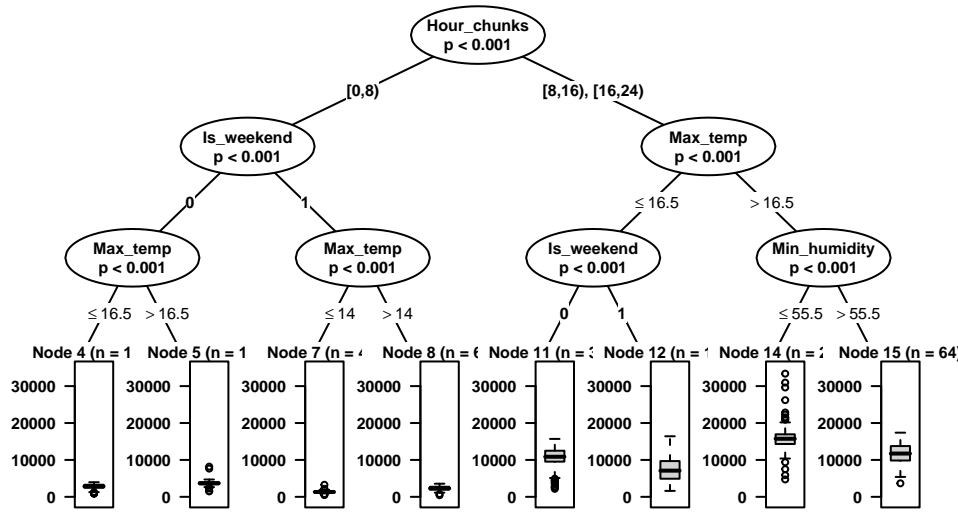


Figure 3: Three layer conditional inference tree

A random forest model was then applied to predict the number of bikes that would be rented. Random forest is composed of many decision trees as opposed to a single tree, resulting in a more accurate result. Incorporating the randomness allows random forest to protect against overfitting and can be applied more effectively to other data sets. The disadvantage of the random forest is its computational complexity. To fit the random forest, we created a method `train_random_forest()` within our `bikeSharing` R package that leverages the `train()` function within the `caret` R package. The optimal tuning parameter `mtry` was determined using a 5-fold cross validation. We visualized a random selected tree.

```
rf_fit <- train_random_forest(data = london_train)
```

Random forest sample trees share similar splits with conditional inference trees (Figure 4). The estimated bike rental count is much higher in the hour chunks that are not 0-8 am. Additionally, the estimated count is high when the temperature is more than 22.25 degrees.

Variable Importance in machine learning models

To investigate which variables affect the prediction most, we calculate the importance of the variables in the conditional inference tree and the random forest. In conditional inference tree, we calculated the mean decrease in accuracy when deleting a variable (Table 1).




Figure 4: A random selected tree in the trained random forest model

Table 1: Mean decrease in accuracy of variables in the conditional inference tree

Variable	Mean decrease in accuracy
Hour_chunks	42292743.475
Max_temp	7305786.627
Is_weekend	2819028.041
Rain_or_snow	1622451.011
Min_humidity	1014266.239
Season	905987.474
Max_humidity	184868.121
Wind_speed	35027.139
Is_holiday	25815.318
Date	3737.395

In random forest, we calculate the increase in MSE when deleting a variable (Figure 5). The code to do this was implemented as a function `plot_rf_importance` in the `bikeSharing` R package.


```
plot_rf_importance(london_train)
```




Figure 5: Feature importance in the random forest model

The top important variables of the conditional inference tree and random forest are similar. The hour chunk is the most important variable in both models, which agrees with the results from the negative binomial GLMM method. For both models, the next most important variable is the Maximum temperature, which is different from the negative GLMM model. Additional details of the variables are discussed in the discussion section. We further compare the performance of random forest and conditional inference tree using the second year bike renting data in London. Conditional inference tree results in a $R^2 = 0.81$ and random forest has $R^2 = 0.91$. Due to the similarity in these two models and the better fit of the random forest model to the training set, we selected random forest tree for use in further comparison.

Methods for Model Fit Assessment

The accuracy of the random forest model and the negative binomial GLMM model will be assessed by using the models to predict the values of bike counts for the London (training set), London (test set), Seoul, and DC data sets. The predicted values will be compared to the actual values to determine how well the models predict the bike counts for a city on a given day and time.

Three metrics of accuracy will be calculated: Root Mean Squared Error (RMSE), Mean Absolute Error (MAE), and Coefficient of Determination (R^2). They will be calculated using the following formulas.

- $RMSE = \sqrt{\frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2}$
- $MAE = \frac{1}{N} \sum_{i=1}^N |y_i - \hat{y}_i|$
- $R^2 = 1 - \frac{\sum_{i=1}^N (y_i - \hat{y}_i)^2}{\sum_{i=1}^N (y_i - \bar{y})^2}$

Results

Both the negative binomial GLMM and the random forest model were used to predict the the values of bike counts for for the London (training set), London (test set), Seoul, and DC data sets. Since Seoul and DC have a different population size than London, it is possible that the predicted value based on the trained London data set and the actual city bike count may be on different scale. In order to account for this, the predicted values for the Seoul and DC data sets were scaled by $\frac{\text{Average Bike Count of City}}{\text{Average Bike Count of London}}$, so that the predicted value for DC and Seoul would be on the same scale as the actual observed bike counts.

The assessment of model fit for the GLMM model was wrapped into the function `glmm_model_fit()`. It takes as input a negative binomial model `model_glmm` which is the output of the `MCEM_algorithm` function, `data`, the data set for which bike counts will be predicted, `scale_to_reference_mean` to determine whether the predicted values should be scaled, and `reference` which is the city that the model was fit to. The output of this function is a vector of RMSE, MAE, and R^2 values. We apply this function to the London training, the London test, the DC data, and the Seoul data.

```
glmm_model_fit(model_glmm = glmm_fit, data = london_train, scale_to_reference_mean = "no",
               reference = london)
```

```
##          RMSE          MAE          R2
## 1 1886.267 1291.106 0.8831618
```

```
glmm_model_fit(model_glmm = glmm_fit, data = london_test, scale_to_reference_mean = "no",
               reference = london)
```

```
##          RMSE          MAE          R2
## 1 2491.293 1647.064 0.8142036
```

```
glmm_model_fit(model_glmm = glmm_fit, data = dc, scale_to_reference_mean = "yes",
               reference = london)
```

```
##          RMSE          MAE          R2
## 1 845.741 605.217 0.521788
```

```
glmm_model_fit(glmm_fit, data = seoul, scale_to_reference_mean = "yes",
               reference = london)
```

```
##          RMSE          MAE          R2
## 1 3519.413 2719.021 0.4999935
```

Unsurprisingly, we see that the GLMM model fits the London test data set the best out of the test sets. There does not seem to be a large difference between the R^2 values for the DC and Seoul data set, however the DC data set has a smaller RMSE and MAE, implying that the DC bike counts may be less spread out than the Seoul bike counts.

A similar assessment for the fit of the random forests model was wrapped into the function `rf_model_fit()`. It takes as input a random forest model, `model`, which is the output of the `train_random_forest` function, `data`, a data set for which bike counts will be predicted applied to, `scale_to_reference_mean` to determine whether the predicted values should be scaled, and `reference` which is the city that the model was fit to. The output of this function is a vector of RMSE, MAE, and R^2 values. We apply this function to the London training, the London test, the DC data, and the Seoul data.

```
rf_model_fit(model = rf_fit, data = london_train, scale_to_reference_mean = "no",
             reference = london)
```

```
##          RMSE          MAE          R2
## 1 752.2336 471.5412 0.9822717
```

```
rf_model_fit(rf_fit, data = london_test, scale_to_reference_mean = "no",
             reference = london)
```

```
##          RMSE          MAE          R2
## 1 1789.24 1191.32 0.9084533
```

```
rf_model_fit(rf_fit, data = dc, scale_to_reference_mean = "yes",
             reference = london)
```

```
##          RMSE          MAE          R2
## 1 664.5474 492.4647 0.6737107
```

```
rf_model_fit(rf_fit, data = seoul, scale_to_reference_mean = "yes",
             reference = london)
```

```
##          RMSE          MAE          R2
## 1 3163.613 2479.835 0.5658297
```


Overall, the random forest model performed better than the negative binomial GLMM in all cases. The random forest model fits the initial training data significantly better than the negative binomial GLMM does. This is not surprising as random forests can account for interactions between covariates while the GLMM model does not. Additionally, the random forest model fits the London test data set better than the GLMM model. Overall, both models seem to perform relatively well in predicting London bike rental counts of the next year (the London training dataset). The random forest model performs significantly better than the GLMM model for DC, however both models perform relatively poorly for the Seoul data set. Thus, it seems that both the GLMM model and the random forest model can predict bike counts in the same city as the testing data set relatively well, however both models do not extrapolate well to other cities.

To determine why this is the case, the average of the covariates for each of the cities was examined to see if there was a major difference between the covariates between cities. (Table 2 below)

Table 2: Average of Covariates From All Data Sets

City	Average Number of Holidays	Average Min Temp	Average Max Temp	Average Min Humidity	Average Max Humidity	Average Wind Speed	Average Rain/Snow Days
London - Train	1.021918	9.576712	15.63836	55.44795	86.58356	4.636123	1.346119
London - Test	1.022018	9.413303	15.44908	56.60183	87.92477	4.227260	1.321101
DC	1.028807	11.384115	19.63523	43.61088	82.14952	3.545157	1.222222
Seoul	1.048159	8.831445	17.45694	37.41643	77.88385	1.726357	1.205855

As we can see, there does not appear to be any major differences in the average of the covariates between the cities. This implies that the issue in applying the fitted models to the new cities is not the range of the covariates, such as the temperature range, humidity range, wind speed range, the number of holidays or the number of rain/snow days. Therefore, to further determine the cause of the poor fit of the model in DC and Seoul, the distribution of the bike counts is examined for each of the data sets below. The DC bike counts and the seoul bike counts were scaled by $\frac{\text{Average Bike Count of London}}{\text{Average Bike Count of City}}$. This was done so that the distribution of bike counts would be on the same scale as the London data set, so that the distributions could be compared more directly.



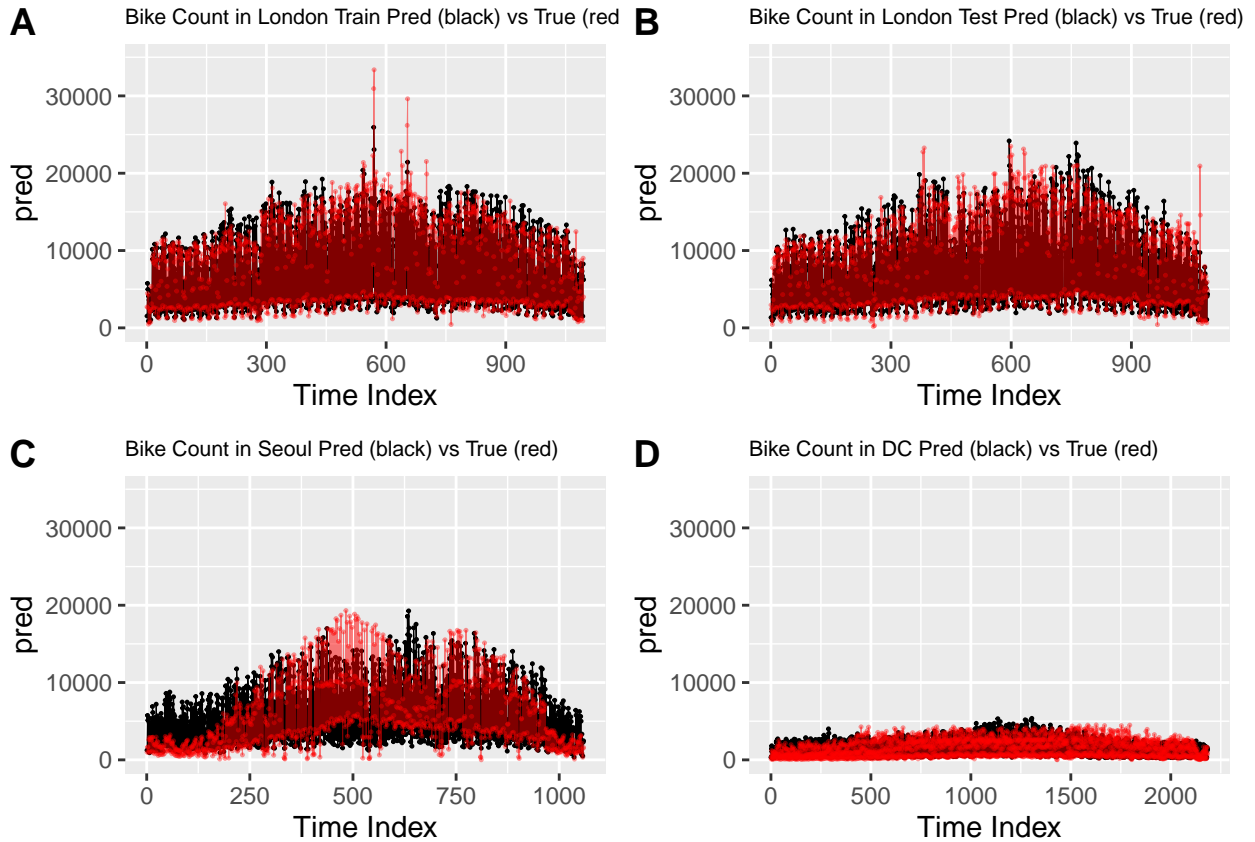
The London training and test set have a very similar distribution, which explains why both the random forest model and the negative binomial glmm model predict the bike counts well for the London test data set. The DC bike counts have a more similar distribution to the London data set than the Seoul bike counts do, which is likely the reason why the random forest model performs better for the DC data compared to the Seoul data. However, overall the distributions between the Seoul bike counts and the DC bike counts are not too different from the London data set. This implies that the effects that the covariates have on bike counts are likely different from city to city. For example, in London, more people ride bikes in the Summer than in the autumn, likely due to the fact that it is warmer. However it is possible that in Seoul or DC, the relationship is the opposite and more people ride bikes in Autumn compared to the Summer because it is too hot in the summer to ride a bike. Thus, since it seems likely that each covariate may influence bike counts differently in different cities, which means our model is not particularly useful in predicting bike counts in cities other than the one that it is trained on. This may be a result of the fact that our cities are very different from each other since they are all from different countries. Further work in measuring bike counts in different cities within a country needs to be done to determine how well our model performs for different cities within the same country as the city that the model was trained on.

Discussion and Conclusion

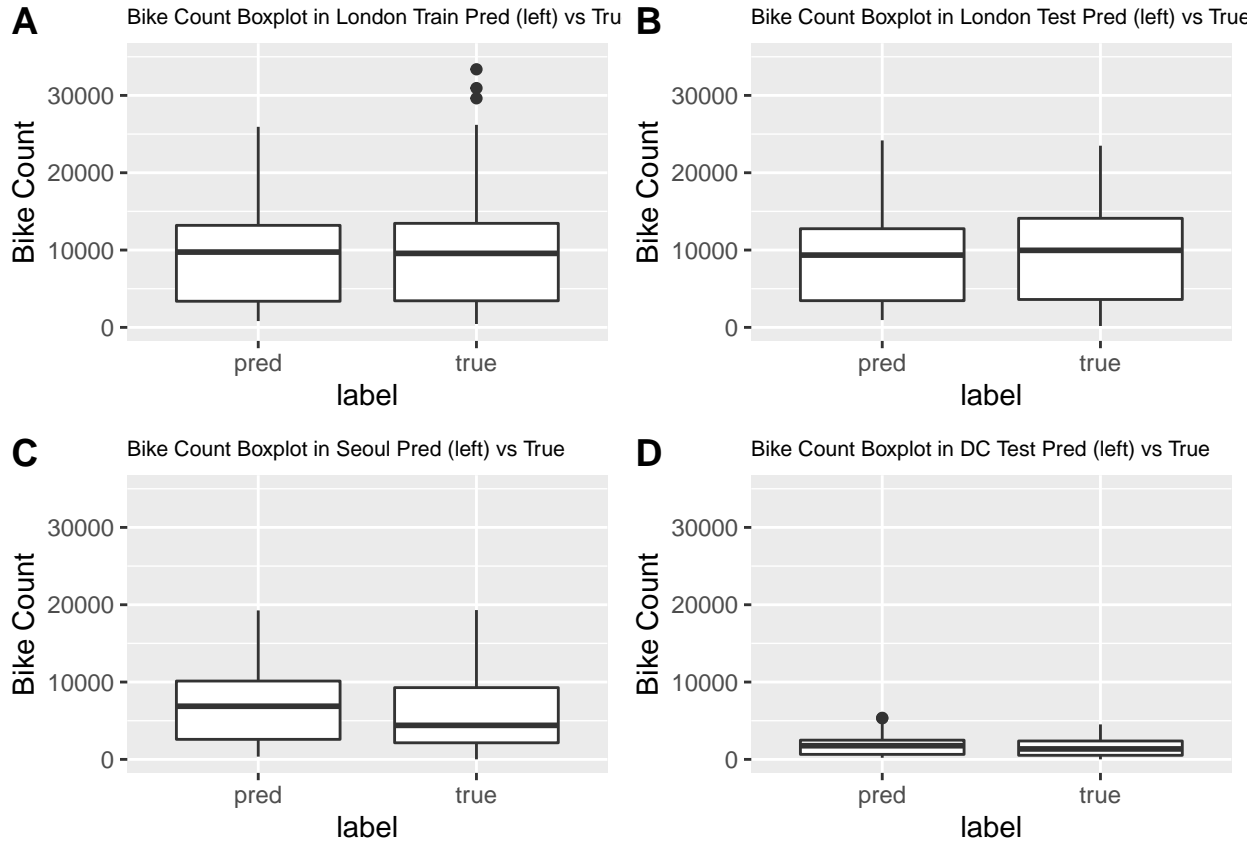
In conclusion, factors including hours, holidays, min temperature, max temperature, min humidity, max humidity, average wind speed and presence of rain/snow all have inference on the number of bike rented in London. Hours from 4pm to midnight have the largest effect on the number of bike rented. Hours from 8am to 4pm, winter season, rain/snow also have a large impact on the bike rent count. Random forest has a better prediction results in all the training and testing data, compared to the results of GLMM.

In addition to the conclusions above, we also have some findings that need more discussion. The RMSE for Seoul is much larger than that for the other two cities. The results for the non London test datasets are worse than the results for either of the two London data sets. This may imply that our model may have some limitations across cities. One possible reason is that in our model, we only considered the effects of days, hours and weather on the bike count but did not take information about cities into consideration (e.g. population, GDP), which may cause bias in prediction.

To get a better understanding on how our negative binomial GLMM model predicts the sequences of bike count in different dataset, we plotted out the predicted bike count and the ground truth for each of the four datasets. Below is the longitudinal plots of the predicted bike count (in black) and the true bike count (in red) for the three cities. As we can see, for Seoul, our models have a tendency to underestimate the predicted bike count in the earlier and later time points, however the model overestimates the bike counts in the middle time points. This is likely why the model has such a larger RMSE and a poor overall fit.



To understand the bias and variance of our prediction results, we also drew the boxplots of predicted bike count and the true bike count of the four datasets. The plots are shown below. As we can see, DC has a very small range of both predicted and true bike counts, which explains why its RMSE and MAE values are much lower than that of the other data sets.



Overall, our result is consistent with the published research by Sylwia et al. (2021) in “Impact of environment on bicycle travel demand-Assessment using bikeshare system data”. In their study, they used the data from the bike sharing system in Cracow, Poland and conducted an ordinal least square regression model to analyze the effect of daily air temperature, daily rainfall, public holidays, and school holidays on the daily number of bike rented from bike sharing system. Their study result indicated that weather conditions, especially air temperature and daily rainfall, have large impact on the number of bike sharing system, which is consistent with our result. Compared to their research, our study further found that the maximum temperature and the minimum humidity have more impact on the bike count. We did not restrict our study on daily level but cut one day into different hour chunks and found that different hour chunks in one day can also have large effect on the number of bike rented from bike sharing system.

Reference

- Hothorn T, Hornik K, Zeileis A (2006). “Unbiased Recursive Partitioning: A Conditional Inference Framework.” *Journal of Computational and Graphical Statistics*, 15(3), 651–674. doi:10.1198/106186006X133933.
- Sylwia P., Mariusz K., Carmelo D (2021). “Impact of environment on bicycle travel demand—Assessment using bikeshare system data.” *Sustainable Cities and Society*, 67, [102724]. <https://doi.org/10.1016/j.scs.2021.102724>
- United Nations, Department of Economic and Social Affairs, Population Division (2018). *World Urbanization Prospects: The 2018 Revision*, Online Edition.